

AOS Assignment 2 Report

Adding Kernel (version: 4.19.210) to the system:

1. Downloading the kernel using the following command:

```
wget
```

```
https://www.kernel.org/pub/linux/kernel/v4.x/linux-4.19.210.tar.xz
```

2. Extracting the kernel source code using:

```
sudo tar -xvf linux-4.17.4.tar.xz -C/usr/src/
```

Installing and Loading the new kernel:

1. After extracting, go to the directory : /usr/src/linux-4.17.4/
2. Compile the kernel: `sudo make`
3. Installing kernel: `sudo make modules_install install`
4. Reboot the system and during the boot you can find a menu where you need to choose Advanced ubuntu settings, followed by a menu where you can choose the kernel version you want to load i.e 4.19.210

Adding system call to the kernel:

1. Go to: /usr/src/linux-4.17.4/
2. Make a directory and create a new C file where you write your new system call code.
3. Inside the system call add necessary headers to create a system call.
4. To create a system call there are multiple ways to create and one of them is

```
SYSCALL_DEFINE# (<name of the system call>)
```

In the above code # indicates the number of arguments you want to add to the function.

5. After writing the system call, we need to add system call entry in `syscall_64.tbl` which is in: `arch/x86/entry/syscalls/`

6. Every system call function should be declared in a header file called `syscalls.h` at the end of the file.

Ex: `asmlinkage long sys_srikanthello(void);`

7. Each system call's object file statement should be written in a Makefile, and the makefile should be in the same folder as the system call.

Ex: `obj-y := hello.o` should be in the makefile.

8. Make file and the c file which contains system calls should be in the same directory.

Ex : If `hello.c` contains a system call which is in a directory `sys_dir`, then the `Makefile` should also be in the same directory.

9. In `/usr/src/linux-4.17.4/` directory , there is a `Makefile` and in this make file we need to include the directory of the file which contains our system call.

Ex: If our system call is defined in `hello.c` which is in `sys_dir` then we need to include `sys_dir` in this make file as `core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/ sys_dir/`

Describing Makefile, sys_64.tbl, syscalls.h :

sys_64.tbl:

```

314  common sched_setattr      __x64_sys_sched_setattr
315  common sched_getattr      __x64_sys_sched_getattr
316  common renameat2         __x64_sys_renameat2
317  common seccomp           __x64_sys_seccomp
318  common getrandom          __x64_sys_getrandom
319  common memfd_create       __x64_sys_memfd_create
320  common kexec_file_load   __x64_sys_kexec_file_load
321  common bpf               __x64_sys_bpf
322  64  execveat            __x64_sys_execveat/ptregs
323  common userfaultfd       __x64_sys_userfaultfd
324  common membarrier        __x64_sys_membarrier
325  common mlock2            __x64_sys_mlock2
326  common copy_file_range   __x64_sys_copy_file_range
327  64  preadyv             __x64_sys_preadv2
328  64  pwritev2            __x64_sys_pwritev2
329  common pkey_mprotect     __x64_sys_pkey_mprotect
330  common pkey_alloc         __x64_sys_pkey_alloc
331  common pkey_free          __x64_sys_pkey_free
332  common statx              __x64_sys_statx
333  common io_pgetevents     __x64_sys_io_pgetevents
334  common rseq               __x64_sys_rseq
335  common q1                 __x64_sys_srikanthello
336  common q2                 __x64_sys_srikantprint
337  common q3                 __x64_sys_srikantprocess
338  common q4                 __x64_sys_srikantgetpid

#
# x32-specific system call numbers start at 512 to avoid cache impact
# for native 64-bit operation. The __x32_compat_sys stubs are created
# on-the-fly for compat_sys_*() compatibility system calls if X86_X32
# is defined.
#
512  x32  rt_sigaction        __x32_compat_sys_rt_sigaction
513  x32  rt_sigreturn         sys32_x32_rt_sigreturn
514  x32  ioctl                __x32_compat_sys_ioctl
515  x32  ready                __x32_compat_sys_ready
516  x32  writev               __x32_compat_sys_writev
517  x32  recvfrom              __x32_compat_sys_recvfrom
518  x32  sendmsg              __x32_compat_sys_sendmsg
519  x32  recvmsg               __x32_compat_sys_recvmsg
520  x32  execve               __x32_compat_sys_execve/ptregs
521  x32  ptrace               __x32_compat_sys_ptrace
522  x32  rt_sigpending         __x32_compat_sys_rt_sigpending
523  x32  rt_sigtimedwait     __x32_compat_sys_rt_sigtimedwait
524  x32  rt_sigqueueinfo      __x32_compat_sys_rt_sigqueueinfo
525  x32  sigaltstack          __x32_compat_sys_sigaltstack
526  x32  timer_create          __x32_compat_sys_timer_create
527  x32  mq_notify             __x32_compat_sys_mq_notify

```

330,1 93%

Above image is *sys_64.tbl*

In the above image, entries 335 to 338 contains the following system calls:

__x64_sys_srikanthello

__x64_sys_srikantprint

__x64_sys_srikantprocess

__x64_sys_srikantgetpid

- First column is the serial number.
- Third column contains the folder name where our system call is present.

- Fourth column contains the name of the system call.

Here we have a prefix `__x64_sys_srikanthello`

`__x64_sys_`: Indicates that this is a system call.

syscalls.h:

```

}

extern int __close_fd(struct files_struct *files, unsigned int fd);

/*
 * In contrast to sys_close(), this stub does not check whether the syscall
 * should or should not be restarted, but returns the raw error codes from
 * __close_fd().
 */
static inline int ksys_close(unsigned int fd)
{
    return __close_fd(current->files, fd);
}

extern long do_sys_open(int dfd, const char __user *filename, int flags,
                      umode_t mode);

static inline long ksys_open(const char __user *filename, int flags,
                           umode_t mode)
{
    if (force_o_largefile())
        flags |= O_LARGEFILE;
    return do_sys_open(AT_FDCWD, filename, flags, mode);
}

extern long do_sys_truncate(const char __user *pathname, loff_t length);

static inline long ksys_truncate(const char __user *pathname, loff_t length)
{
    return do_sys_truncate(pathname, length);
}

static inline unsigned int ksys_personality(unsigned int personality)
{
    unsigned int old = current->personality;

    if (personality != 0xffffffff)
        set_personality(personality);

    return old;
}
asm linkage long sys_hello(void);
asm linkage long sys_srikanthello(void);
asm linkage long sys_srikantprint(char* );
asm linkage long sys_srikantprocess(void);
asm linkage long sys_srikantpgetpid(void);

#endif

```

1301,1 Bot

All the system calls declarations should be added at the end of the file, before the end of last if block i.e `#endif`

Makefile:

```
ifeq ($KBUILD_EXTMOD,,)
      core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ hello/ q1/ q2/ q3/ q4/
```

You can find the above line starting with `core-y` in Makefile which is in line: 987

Add the folders where you defined your system call.

Here `q1/, q2/, q3/, q4/` are folders containing `srikanthello, srikantprint, srikantprocess, srikantgetpid` system calls respectively.

Describing system calls:

1. srikanthello()

```
#include <linux/kernel.h>
#include <linux/syscalls.h>

SYSCALL_DEFINE0(srikanthello)
{
    printk("Hello world\n");
    return 0;
}
```

- In the above image, the system call is defined using `SYSCALL_DEFINE0` where 0 indicates the number of parameters.
- This system call prints **Hello world** inside the log file.
- `printk` function is used to print our messages in the log file.
- Here k indicates: kernel

2. srikantprint()

```
#include <linux/kernel.h>
#include <linux/linkage.h>
#include <linux/syscalls.h>
#include <linux/uaccess.h>

SYSCALL_DEFINE1(srikantprint, char *, msg)
{
    char buffer[256];
    long chunk = strncpy_from_user(buffer, msg, sizeof(buffer));
    if (chunk < 0 || chunk == sizeof(buffer))
        return -EFAULT;
    printk(KERN_INFO "SYSTEM CALL srikantprint: \"%s\"\n", buffer);
    return 0;
}
```

- In the above image, the system call is defined using `SYSCALL_DEFINE1` where 1 indicates the number of parameters.
- This system call takes one string argument and prints it in the log file.

3. srikantprocess()

```
#include <linux/kernel.h>
#include <linux/sched.h>
#include <linux/cred.h>
#include <linux/syscalls.h>

SYSCALL_DEFINE0(srikantprocess)
{
    struct task_struct *parent = current->parent;
    printk("Process ID: %d ; Parent process ID: %d",current->pid,parent->pid);
    return 0;
}
```

- In the above image, the system call is defined using `SYSCALL_DEFINE0` where 0 indicates the number of parameters.
- Here this system call prints the process IDs of parent and the child.
- Predefined struct called `task_struct` is used to get the respective IDs.

4. srikantgetpid()

```
#include <linux/kernel.h>
#include <linux/syscalls.h>
#include <linux/sched.h>
#include <linux/cred.h>

SYSCALL_DEFINE0(srikantgetpid)
{
    return task_tgid_vnr(current);
}
```

- In the above image system call is defined using `SYSCALL_DEFINE0` where 0 indicates the number of parameters.
- This system call returns the process ID of the process which invoked this system call.
- Here `task_tgid_vnr` is the system call which is invoked to get the process ID of the process.

Testing:

```
#define __GNU_SOURCE
#include <unistd.h>
#include <sys/syscall.h>
#include <stdio.h>

int main()
{
    long int amma = syscall(335);
    printf("System call srikanthello returned: %ld\n",amma);

    int res1=syscall(337);
    printf("System call srikantprocess returned: %d \n",res1);

    long int res2=syscall(338);
    printf("System call srikantgetpid returned: %ld\n",res2);

    char *arg = "Srikant Konduri";
    printf("Making system call with \"%s\n\", arg");
    long res = syscall(336,arg);
    printf("System call srikantprint returned: %ld\n", res);
    return 0;
}
```

Output:

```
[azureuser@SkVM:~$ ./a.out
System call srikanthello returned: 0
System call srikantprocess returned: 0
System call srikantgetpid returned: 1276
Making system call with "Srikant Konduri"
System call srikantprint returned: 0
azureuser@SkVM:~$ ]
```

After running the above test file named test.c we got the above returned values.

- System call *srikanthello()* returns 0 on successful call and logs *hello world* in the log file.
- System call *srikantprocess()* returns 0 on successful call and logs process ID of parent and current process in the log file.
- System call *srikantgetpid()* return the process ID of a.out (i.e test.c) which is the current process.
- System call *srikantprint()* prints the argument in the log file which we have sent. Here “Srikant Konduri” is sent as an argument.

```
[ 5.478866] audit: type=1400 audit(1663346158.988:13): apparmor="STATUS" operation="profile_load" profile="unconfined" name="man_groff" pid=518 comm="apparmor_parser"
[ 5.483288] audit: type=1400 audit(1663346158.984:14): apparmor="STATUS" operation="profile_load" profile="unconfined" name="lsb_release" pid=520 comm="apparmor_parser"
[ 5.486477] audit: type=1400 audit(1663346158.988:15): apparmor="STATUS" operation="profile_load" profile="unconfined" name="/usr/sbin/tcpdump" pid=521 comm="apparmor_parser"
[ 5.531992] audit: type=1400 audit(1663346159.032:16): apparmor="STATUS" operation="profile_load" profile="unconfined" name="/snap/snapd/16292/usr/lib/snapd/snap-confine" pid=534 comm="apparmor_parser"
[ 5.531996] audit: type=1400 audit(1663346159.032:17): apparmor="STATUS" operation="profile_load" profile="unconfined" name="/snap/snapd/16292/usr/lib/snapd/snap-confine//mount-namespace-capture-helper" pid=534 comm="apparmor_parser"
[ 5.534266] audit: type=1400 audit(1663346159.036:18): apparmor="STATUS" operation="profile_load" profile="unconfined" name="/snap/snapd/16778/usr/lib/snapd/snap-confine" pid=535 comm="apparmor_parser"
[ 6.174932] mlx5_core: 971e:00:02.0 enP38686s1: Link up
[ 6.175657] hv_netvsc: 6845bdae-92bb-6845-bdae-92bb045bdae eth0: Data path switched to VF: enP38686s1
[ 6.176834] IPv6: ADDRCONF(NETDEV_UP): eth0: link is not ready
[ 6.176851] IPv6: ADDRCONF(NETDEV_CHANGE): eth0: link becomes ready
[ 6.792715] hv_netvsc: 6845bdae-92bb-6845-bdae-92bb045bdae eth0: Data path switched from VF: enP38686s1
[ 8.959447] mlx5_core: 971e:00:02.0 enP38686s1: Link up
[ 8.960369] hv_netvsc: 6845bdae-92bb-6845-bdae-92bb045bdae eth0: Data path switched to VF: enP38686s1
[ 9.136876] hv_netvsc: 6845bdae-92bb-6845-bdae-92bb045bdae eth0: Data path switched from VF: enP38686s1
[ 9.498569] mlx5_core: 971e:00:02.0 enP38686s1: Link up
[ 9.499333] hv_netvsc: 6845bdae-92bb-6845-bdae-92bb045bdae eth0: Data path switched to VF: enP38686s1
[ 15.289759] EXT4-fs (sdb1): mounted filesystem with ordered data mode. Opts: (null)
[ 18.412837] kaudiotd_printk_skbe: 14 callbacks suppressed
[ 18.412839] audit: type=1400 audit(1663346171.912:33): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="/snap/snapd/16778/usr/lib/snapd/snap-confine" pid=978 comm="apparmor_parser"
[ 18.412843] audit: type=1400 audit(1663346171.912:34): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="/snap/snapd/16778/usr/lib/snapd/snap-confine//mount-namespace-capture-helper" pid=978 comm="apparmor_parser"
[ 18.415828] audit: type=1400 audit(1663346171.916:35): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap-update-nx.lxd" comm="apparmor_parser"
[ 18.418481] audit: type=1400 audit(1663346171.920:36): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap.lxd.activate" pid=981 comm="apparmor_parser"
[ 18.418783] audit: type=1400 audit(1663346171.920:37): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap.lxd.benchmark" pid=982 comm="apparmor_parser"
[ 18.421278] audit: type=1400 audit(1663346171.920:38): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap.lxd.buginfo" pid=983 comm="apparmor_parser"
[ 18.421356] audit: type=1400 audit(1663346171.920:39): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap.lxd.check-kernel" pid=984 comm="apparmor_parser"
[ 18.424922] audit: type=1400 audit(1663346171.924:40): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap.lxd.hook.configure" pid=986 comm="apparmor_parser"
[ 18.425162] audit: type=1400 audit(1663346171.924:41): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap.lxd.daemon" pid=985 comm="apparmor_parser"
[ 18.427457] audit: type=1400 audit(1663346171.928:42): apparmor="STATUS" operation="profile_replace" info="same as current profile, skipping" profile="unconfined" name="snap.lxd.hook.install" pid=987 comm="apparmor_parser"
[ 49.184437] hv_balloon: Max. dynamic memory size: 8192 MB
[ 50.213304] Hello world
[ 50.213388] Process ID: 1276 ; Parent process ID: 1211
[ 50.213415] SYSTEM CALL srikantprint: "Srikant Konduri"
azureuser@SkVM:~$ ]
```

Answer to Question 3:

- Here current process ID and parent process ID are different.
- Parent process ID is the process ID of the process which is invoking.
- Here Parent process is 1211 which is the ID of bash which is invoking a.out (test.c)
- Current process ID is 1276 which is the process ID of a.out